

Jiashu Xu 徐家澍

📍 Cambridge, MA | 📞 949-522-2936 | ✉️ jxu1@g.harvard.edu | 🌐 Website | 🎓 Scholar | 📄 [cnut1648](#) | 📱 [jiashu-xu](#)

EDUCATION

Harvard University

M.S. in Computational Science and Engineering; cross-registered at MIT; GPA: 4.0/4.0

Cambridge, USA

Fall 2022 – Spring 2024

University of Southern California

B.S. double major in Applied Math & Computer Science; **Summa Cum Laude**; GPA: 3.97/4.0

Los Angeles, USA

Fall 2020 – Spring 2022

University of California, Irvine

B.S. double major in Applied Math & Computer Science; GPA: 3.98/4.0

Irvine, USA

Fall 2018 – Spring 2020

Hong Kong University of Science and Technology

UCEAP summer visiting student; study robotics; GPA: 4.0/4.0

Hong Kong, China

Summer 2019

AWARDS

CURVE Research Fellowship

\$1250/semester Research Stipend

Jennifer Battat Scholarship

\$3.5k for Mathematics Major

USC Transfer Merit Scholarship

Half-tuition Merit Scholarship for 2% Of Transfer Applicants

USC Academic Achievement Award

Double Major with 3.75+ GPA

USC & UCI Dean's List

All Semesters

RESEARCH INTERESTS

My current research interests lie in **Reliable AI**. Particularly,

1. AI Safety. For example, defending against malicious exploitation of LLM vulnerabilities ([3], [4]), protecting open-sourced LLM ownership ([1]), and aligning LLMs with human preferences ([2]).
2. Training AI that excels in low-resource regimes, through indirect supervision ([7], [11]) or synthetic data generation ([5], [6], [8], [10]).
3. Explanation and model learning from explanation ([12] to [14]).

PUBLICATIONS & SERVICES

***=EQUAL CONTRIBUTION**

[1] Training Large Language Models as Reward Models

Jiashu Xu, Daniel Pressel, Prasoon Goyal, Luke Dai, Reza Ghanadan, Michael Johnston
COLM, 2024 (To be submitted, under Amazon Internal Review)

paper

[2] Instructional Fingerprinting of Large Language Models

Jiashu Xu, Fei Wang*, Mingyu Derek Ma*, Pang Wei Koh, Chaowei Xiao, Muchao Chen
NAACL, 2024 (Under Review)

code paper project

[3] Instructions as Backdoors: Backdoor Vulnerabilities of Instruction Tuning for Large Language Models

Jiashu Xu, Mingyu Derek Ma, Fei Wang, Chaowei Xiao, Muhao Chen
NAACL, 2024 (Under Review)

paper

[4] Test-time Backdoor Mitigation for Black-Box Large Language Models with Defensive Demonstrations

Wenjie Mo, **Jiashu Xu**, Qin Liu, Jiong Xiao Wang, Jun Yan, Chaowei Xiao, Muhao Chen
NAACL, 2024 (Under Review)

paper

[5] BEHAVIOR Vision Suite: Customizable Dataset Generation via Simulation

Yunhao Ge*, Yihe Tang*, **Jiashu Xu***, Cem Gokmen*, Chengshu Li, Wensi Ai, Benjamin Jose Martinez, Arman Aydin, Mona Anvari, Ayush K Chakravarthy, Hong-Xing Yu, Josiah Wong, Sanjana Srivastava, Sharon Lee, Shengxin Zha, Laurent Itti, Yunzhu Li, Roberto Martín-Martín, Miao Liu, Pengchuan Zhang, Ruohan Zhang, Fei-Fei Li, Jiajun Wu
CVPR, 2024 (Under Review)

paper

- [6] **DreamDistribution: Prompt Distribution Learning for Text-to-Image Diffusion Models**
 Brian Nlong Zhao, Yuhang Xiao*, **Jiashu Xu***, Xinyang Jiang, Yifan Yang, Dongsheng Li, Laurent Itti, Yunhao Ge, Vibhav Vineet
CVPR, 2024 (Under Review) code paper project
- [7] **Can NLI Provide Proper Indirect Supervision for Low-resource Biomedical Relation Extraction?**
Jiashu Xu, Mingyu Derek Ma, Muhao Chen
ACL, 2023 (**Oral**) code paper
- [8] **Dall-e for detection: Language-driven context image synthesis for object detection**
 Yunhao Ge*, **Jiashu Xu***, Brian Nlong Zhao, Neel Joshi, Laurent Itti, Vibhav Vineet
arXiv, 2022 code paper extension
- [9] **X-Norm: Exchanging Normalization Parameters for Bimodal Fusion**
 Yufeng Yin*, **Jiashu Xu***, Tianxin Zu, Mohammad Soleymani
ICMI, 2022 paper
- [10] **Neural-Sim: Learning to Generate Training Data with NeRF**
 Yunhao Ge, Harkirat Behl*, **Jiashu Xu***, Suriya Gunasekar, Neel Joshi, Yale Song, Xin Wang, Laurent Itti, Vibhav Vineet
ECCV, 2022 code paper
- [11] **Unified Semantic Typing with Meaningful Label Inference**
 James Y. Huang, Bangzheng Li*, **Jiashu Xu***, Muhao Chen
NAACL, 2022 code paper
- [12] **Dissection Gesture Sequence during Nerve Sparing Predicts Erectile Function Recovery after Robot-Assisted Radical Prostatectomy**
 Runzhuo Ma, **Jiashu Xu**, Ivan Rodriguez, Gina DeMeo, Aditya Desai, Loc Trinh, Jessica H. Nguyen, Anima Anandkumar, Jim C. Hu, Andrew J. Hung
NPJ Digital Medicine, 2022 paper
- [13] **Dissection Assessment for Robotic Technique (DART) to Evaluate Nerve-Spare of Robot-Assisted Radical Prostatectomy**
 Runzhuo Ma, Alvin Hui, **Jiashu Xu**, Aditya Desai, Michael Tzeng, Emily Cheng, Loc Trinh, Jessica H. Nguyen, Anima Anandkumar, Jim C. Hu, Andrew J. Hung
American Urological Association Annual Conference (AUA), 2022 paper
- [14] **SalKG: Learning From Knowledge Graph Explanations for Commonsense Reasoning**
 Aaron Chan, **Jiashu Xu**, Boyuan Long, Soumya Sanyal, Tanishq Gupta, Xiang Ren
NeurIPS, 2021 code paper

Reviewer Service: ACL Rolling Review, ACL 2023, EMNLP 2023, CVPR 2022

RESEARCH EXPERIENCE

NVIDIA Research

Santa Clara, USA

(Incoming) Research Scientist Intern | Manager: [Ming-Yu Liu](#)

Summer 2024

- Plan to work on LLM.

Amazon Science

New York, USA

Applied Scientist Intern | Manager: [Daniel Pressel](#), [Michael Johnston](#)

Summer 2023

- Collaborated closely with the LLM team on the reward modeling side.
- Finetuned LLMs directly as reward models such that models learn to align with human preferences implicitly. Further benefits included zero-shot generalization to unseen dimensions and domains, high-quality data filtering, rationale generation to explain decisions, and synthetic conversation curation for AI self-improvement (RLAIF) [1].

USC LUKA Group

Los Angeles, USA

Research Assistant | Advisor: Prof. [Muhao Chen](#), Prof. [Chaowei Xiao](#)

Fall 2021 – Present

- Proposed a fingerprinting method to safeguard open-source LLM ownership via memorizing fingerprint instances. Such lightweight fingerprint persists large-scale user finetuning on arbitrary datasets, is robust to fingerprint guessing and various PEFT training methods, and supports multi-stage fingerprinting akin to MIT License [2].

- Investigated backdoor vulnerabilities of instruction-tuned LLMs that have high backdoor success with minimal malicious instructions, can generalize to multiple tasks, and cannot be cured by continual learning [3]. And proposed leveraging clean in-context demonstrations as effective test-time defense against various backdoor attacks [4].
- Proposed indirect supervision to borrow supervision signals from resource-rich tasks to enhance resource-limited tasks: cross-domain transfer general domain NLI knowledge to improve low-resource biomedical Relation Extraction [7]; cross-task transfer semantic typing knowledge to handle large label space inference [11].

Harvard AI4LIFE Group

Cambridge, USA

Research Assistant | Advisor: Prof. [Himabindu Lakkaraju](#)

Spring 2023 – Present

- Integrated dynamic knowledge graph, constructed with a language model agent dynamically, into inference to improve factuality (In Progress).
- Investigated mechanistic interpretability on LLaVA-like vision language models to investigate how models react to complex queries such as referring expressions and object counting (In Progress).

Stanford SVL Group

Palo Alto, USA

Research Assistant | Advisor: Prof. [Jiajun Wu](#)

Fall 2023 – Present

- Developed BEHAVIOR Vision Suite, a customizable dataset generator featuring photorealistic assets and physically plausible annotations. Demonstrated applications include holistic benchmarks for 2D and 3D vision models, robustness evaluation through parametric out-of-distribution evaluation (*e.g.* low lighting, extreme camera pose), and synthetic dataset generation to bolster performance in low-resource scenarios [5].

Microsoft Research

Los Angeles, USA

Student Collaborator | Manager: Prof. [Laurent Itti](#), [Vibhav Vineet](#)

Spring 2022 – Present

- Proposed prompt distribution learning for text-to-image and text-to-3D diffusion models to lightweight control image quality and diversity [6].
- Utilized diffusion models and object cut-and-paste to create coherent synthetic training datasets for enhancing low-resource object detection and segmentation [8]. And proposed differentiable synthetic dataset generation with NeRF to improve out-of-distribution object detection of varying views [10].

MENTORING

[Wenjie Mo](#) USC B.S.

Fall 2023 – Present

[Brian Nlong Zhao](#) USC B.S. → M.S., Research Scientist Intern at Microsoft Research Asia

Fall 2022 – Fall 2023

WORKING & TEACHING EXPERIENCE

Teaching Assistant at USC

Los Angeles, USA

Role: Office Hours, Discussion Sections, Grading

Spring – Fall 2021

- CSCI 567: Graduate level Machine Learning with Prof. [Haipeng Luo](#).
- MATH 499: Independent Research with Prof. [Neelesh Tiruvilumala](#).

Teach for Los Angeles

Los Angeles, USA

Mentor

Spring 2021

- Tutored middle school students from LA K-12 community 1-on-1 on mathematics for two hours every week.
- Inspired students to reach full math potential in preparation for college and STEM careers.

Math CEO

Irvine, USA

Mentor

Fall 2018 – Spring 2020

- Coordinated meetings with Santa Ana middle school students and taught mathematical thinking.

Johnson & Johnson

Shanghai, China

Digital & Analytics Data Assistant

Summer 2019

- Tracked counterfeit or parallel products from various sales channels using NLP techniques including semantic role labeling and named entity recognition.
- Devised a context extractor based on Jieba tokenizer and Chinese word vectors.
- Presented at PCS 2019 medicine CIO summit about the NLP approach for tracking counterfeit products.

Wind Information

Shanghai, China

Quantitative Index Research Analyst

Spring – Summer 2018

- Collaborated with product managers to launch Wind's new product: Wind Equity Backtester and implemented multiple prototype algorithms with test codes using python-wind and Pytest.
- Code-reviewed index-related codes and queried Wind index database to resolve clients' complaints.