Jiashu Xu 徐家澍

📍 Cambridge, MA | 📱 949-522-2936 | ✉ jxu1@g.harvard.edu | 🌐 Website | 🎓 Scholar | 🐙 cnut1648 | 💼 jiashu-xu

EDUCATION

Harvard University — Cambridge, USA
M.S. in Computational Science and Engineering; cross-registered at MIT; GPA: 4.0/4.0 — *Fall 2022 – Spring 2024*

University of Southern California — Los Angeles, USA
B.S. double major in Applied Math & Computer Science; Summa Cum Laude; GPA: 3.97/4.0 — *Fall 2020 – Spring 2022*

University of California, Irvine — Irvine, USA
B.S. double major in Applied Math & Computer Science; GPA: 3.98/4.0 — *Fall 2018 – Spring 2020*

Hong Kong University of Science and Technology — Hong Kong, China
UCEAP summer visiting student; study robotics; GPA: 4.0/4.0 — *Summer 2019*

AWARDS

CURVE Research Fellowship	$1250/semester Research Stipend
Jennifer Battat Scholarship	$3.5k for Mathematics Major
USC Transfer Merit Scholarship	Half-tuition Merit Scholarship for 2% Of Transfer Applicants
USC Academic Achievement Award	Double Major with 3.75+ GPA
USC & UCI Dean's List	All Semesters

RESEARCH INTERESTS

My current research interests lie in **Reliable AI**. Particularly,

1. AI Safety. For example, defending against malicious exploitation of LLM vulnerabilities ([3], [4]), protecting open-sourced LLM ownership ([1]), and aligning LLMs with human preferences ([2]).
2. Training AI that excels in low-resource regimes, through indirect supervision ([7], [11]) or synthetic data generation ([5], [6], [8], [10]).
3. Explanation and model learning from explanation ([12] to [14]).

PUBLICATIONS & SERVICES   *=EQUAL CONTRIBUTION

[1] **Training Large Language Models as Reward Models**
Jiashu Xu, Daniel Pressel, Prasoon Goyal, Luke Dai, Reza Ghanadan, Michael Johnston
COLM, 2024 (To be submitted, under Amazon Internal Review) — paper

[2] **Instructional Fingerprinting of Large Language Models**
Jiashu Xu, Fei Wang*, Mingyu Derek Ma*, Pang Wei Koh, Chaowei Xiao, Muchao Chen
NAACL, 2024 — code paper project

[3] **Instructions as Backdoors: Backdoor Vulnerabilities of Instruction Tuning for Large Language Models**
Jiashu Xu, Mingyu Derek Ma, Fei Wang, Chaowei Xiao, Muhao Chen
NAACL, 2024 — paper

[4] **Test-time Backdoor Mitigation for Black-Box Large Language Models with Defensive Demonstrations**
Wenjie Mo, **Jiashu Xu**, Qin Liu, Jiongxiao Wang, Jun Yan, Chaowei Xiao, Muhao Chen
COLM, 2024 (To be submitted) — paper

[5] **BEHAVIOR Vision Suite: Customizable Dataset Generation via Simulation**
Yunhao Ge*, Yihe Tang*, **Jiashu Xu**, Cem Gokmen*, Chengshu Li, Wensi Ai, Benjamin Jose Martinez, Arman Aydin, Mona Anvari, Ayush K Chakravarthy, Hong-Xing Yu, Josiah Wong, Sanjana Srivastava, Sharon Lee, Shengxin Zha, Laurent Itti, Yunzhu Li, Roberto Martín-Martín, Miao Liu, Pengchuan Zhang, Ruohan Zhang, Fei-Fei Li, Jiajun Wu
CVPR, 2024 — paper

[6] **DreamDistribution: Prompt Distribution Learning for Text-to-Image Diffusion Models**
Brian Nlong Zhao, Yuhang Xiao*, **Jiashu Xu***, Xinyang Jiang, Yifan Yang, Dongsheng Li, Laurent Itti, Yunhao Ge, Vibhav Vineet
ECCV, 2024 (Under Review)   `code` `paper` `project`

[7] **Can NLI Provide Proper Indirect Supervision for Low-resource Biomedical Relation Extraction?**
Jiashu Xu, Mingyu Derek Ma, Muhao Chen
ACL, 2023 (**Oral**)   `code` `paper`

[8] **Dall-e for detection: Language-driven context image synthesis for object detection**
Yunhao Ge*, **Jiashu Xu***, Brian Nlong Zhao, Neel Joshi, Laurent Itti, Vibhav Vineet
arXiv, 2022   `code` `paper` `extension`

[9] **X-Norm: Exchanging Normalization Parameters for Bimodal Fusion**
Yufeng Yin*, **Jiashu Xu***, Tianxin Zu, Mohammad Soleymani
ICMI, 2022   `paper`

[10] **Neural-Sim: Learning to Generate Training Data with NeRF**
Yunhao Ge, Harkirat Behl*, **Jiashu Xu***, Suriya Gunasekar, Neel Joshi, Yale Song, Xin Wang, Laurent Itti, Vibhav Vineet
ECCV, 2022   `code` `paper`

[11] **Unified Semantic Typing with Meaningful Label Inference**
James Y. Huang, Bangzheng Li*, **Jiashu Xu***, Muhao Chen
NAACL, 2022   `code` `paper`

[12] **Dissection Gesture Sequence during Nerve Sparing Predicts Erectile Function Recovery after Robot-Assisted Radical Prostatectomy**
Runzhuo Ma, **Jiashu Xu**, Ivan Rodriguez, Gina DeMeo, Aditya Desai, Loc Trinh, Jessica H. Nguyen, Anima Anandkumar, Jim C. Hu, Andrew J. Hung
NPJ Digital Medicine, 2022   `paper`

[13] **Dissection Assessment for Robotic Technique (DART) to Evaluate Nerve-Spare of Robot-Assisted Radical Prostatectomy**
Runzhuo Ma, Alvin Hui, **Jiashu Xu**, Aditya Desai, Michael Tzeng, Emily Cheng, Loc Trinh, Jessica H. Nguyen, Anima Anandkumar, Jim C. Hu, Andrew J. Hung
American Urological Association Annual Conference (AUA), 2022   `paper`

[14] **SalKG: Learning From Knowledge Graph Explanations for Commonsense Reasoning**
Aaron Chan, **Jiashu Xu**, Boyuan Long, Soumya Sanyal, Tanishq Gupta, Xiang Ren
NeurIPS, 2021   `code` `paper`

Reviewer Service: ACL Rolling Review, ACL 2023, EMNLP 2023, CVPR 2022

Research Experience

NVIDIA Research   Santa Clara, USA
(Incoming) Research Scientist Intern | Manager: Ming-Yu Liu   *Summer 2024*
- Plan to work on LLM.

Amazon Science   New York, USA
Applied Scientist Intern | Manager: Daniel Pressel, Michael Johnston   *Summer 2023*
- Collaborated closely with the LLM team on the reward modeling side.
- Finetuned LLMs directly as reward models such that models learn to align with human preferences implicitly. Further benefits included zero-shot generalization to unseen dimensions and domains, high-quality data filtering, rationale generation to explain decisions, and synthetic conversation curation for AI self-improvement (RLAIF) [1].

USC LUKA Group   Los Angeles, USA
Research Assistant | Advisor: Prof. Muhao Chen, Prof. Chaowei Xiao   *Fall 2021 − Present*
- Proposed a fingerprinting method to safeguard open-source LLM ownership via memorizing fingerprint instances. Such lightweight fingerprint persists large-scale user finetuning on arbitrary datasets, is robust to fingerprint guessing and various PEFT training methods, and supports multi-stage fingerprinting akin to MIT License [2].

- Investigated backdoor vulnerabilities of instruction-tuned LLMs that have high backdoor success with minimal malicious instructions, can generalize to multiple tasks, and cannot be cured by continual learning [3]. And proposed leveraging clean in-context demonstrations as effective test-time defense against various backdoor attacks [4].
- Proposed indirect supervision to borrow supervision signals from resource-rich tasks to enhance resource-limited tasks: cross-domain transfer general domain NLI knowledge to improve low-resource biomedical Relation Extraction [7]; cross-task transfer semantic typing knowledge to handle large label space inference [11].

Harvard AI4LIFE Group
Research Assistant | Advisor: Prof. Himabindu Lakkaraju

Cambridge, USA
Spring 2023 − Present
- Integrated dynamic knowledge graph, constructed with a language model agent dynamically, into inference to improve factuality (In Progress).
- Investigated mechanistic interpretability on LLaVA-like vision language models to investigate how models react to complex queries such as referring expressions and object counting (In Progress).

Stanford SVL Group
Research Assistant | Advisor: Prof. Jiajun Wu

Palo Alto, USA
Fall 2023 − Present
- Developed BEHAVIOR Vision Suite, a customizable dataset generator featuring photorealistic assets and physically plausible annotations. Demonstrated applications include holistic benchmarks for 2D and 3D vision models, robustness evaluation through parametric out-of-distribution evaluation (*e.g.* low lighting, extreme camera pose), and synthetic dataset generation to bolster performance in low-resource scenarios [5].

Microsoft Research
Student Collaborator | Manager: Prof. Laurent Itti, Vibhav Vineet

Los Angeles, USA
Spring 2022 − Present
- Proposed prompt distribution learning for text-to-image and text-to-3D diffusion models to lightweight control image quality and diversity [6].
- Utilized diffusion models and object cut-and-paste to create coherent synthetic training datasets for enhancing low-resource object detection and segmentation [8]. And proposed differentiable synthetic dataset generation with NeRF to improve out-of-distribution object detection of varying views [10].

Mentoring

Wenjie Mo USC B.S.	Fall 2023 − Present
Brian Nlong Zhao USC B.S. → M.S., Research Scientist Intern at Microsoft Research Asia	Fall 2022 − Fall 2023

Working & Teaching Experience

Teaching Assistant at USC
Role: Office Hours, Discussion Sections, Grading

Los Angeles, USA
Spring − Fall 2021
- CSCI 567: Graduate level Machine Learning with Prof. Haipeng Luo.
- MATH 499: Independent Research with Prof. Neelesh Tiruviluamala.

Teach for Los Angeles
Mentor

Los Angeles, USA
Spring 2021
- Tutored middle school students from LA K-12 community 1-on-1 on mathematics for two hours every week.
- Inspired students to reach full math potential in preparation for college and STEM careers.

Math CEO
Mentor

Irvine, USA
Fall 2018 − Spring 2020
- Coordinated meetings with Santa Ana middle school students and taught mathematical thinking.

Johnson & Johnson
Digital & Analytics Data Assistant

Shanghai, China
Summer 2019
- Tracked counterfeit or parallel products from various sales channels using NLP techniques including semantic role labeling and named entity recognition.
- Devised a context extractor based on Jieba tokenizer and Chinese word vectors.
- Presented at PCS 2019 medicine CIO summit about the NLP approach for tracking counterfeit products.

Wind Information
Quantitative Index Research Analyst

Shanghai, China
Spring − Summer 2018
- Collaborated with product managers to launch Wind's new product: Wind Equity Backtester and implemented multiple prototype algorithms with test codes using python-wind and Pytest.
- Code-reviewed index-related codes and queried Wind index database to resolve clients' complaints.